Sanjana Kulshreshtha

Data Engineer

CONTACT

sanjana.kulshreshtha@example.com
Pune, India

SKILLS

- PySpark
- AWS EMR
- Kinesis
- Airflow
- Python
- Hive
- Presto
- SQL
- pytest
- Parquet
- Redshift Spectrum
- Helm charts for job deployment
- Airflow plugins

LANGUAGES

- English (fluent)
- Hindi (native)
- Marathi (native)

EDUCATION

- BE Computer Engineering, Sahyadri College of Engineering (2013 – 2017)

SUMMARY

Data engineer with seven years in e-commerce, mostly building ingestion frameworks and reusable libraries. I like config-driven pipelines and small, sharp tools. Wrote an internal Spark testing helper now used by three teams. I mentor new joiners and run a monthly internal talk about data quality practices. I enjoy simplifying complicated workflows and writing documentation that new colleagues actually read.

EXPERIENCE

Senior Data Engineer, Bazaarnest Commerce 2020 – Present

- Designed a config-driven ingestion framework: onboarding a new source went from two weeks to two days.
- Processed 800 million clickstream events daily on EMR with Spark Structured Streaming.
- Built a PySpark unit testing library with fixtures for common transformations.
- Deployed batch jobs on the shared platform's Kubernetes cluster using Helm charts supplied by the platform team.
- Collaborated with data scientists to deliver feature tables for recommendation models, with documented freshness guarantees.

Data Engineer, Mandi Fresh Foods 2017 – 2020

- Loaded supplier feeds into Hive tables.
- Optimised slow joins and skewed partitions.
- Created a metadata catalogue so teams could find and trust supplier datasets.
- Helped migrate on-premise Hive tables to cloud storage in under four months.
- Cut daily supplier loads from three hours to forty minutes through partition tuning and incremental processing.

Software Engineer, Konkan Infosystems 2015 – 2017

- Developed Java web services and SQL reports for a logistics client.
- Fixed production defects and wrote unit tests.
- Learned agile practices in a team of twelve developers.
- Added data quality checks to critical tables and alerted owners automatically.
- Supported ad hoc analyst requests with fast, clean extracts.